

Assignment Project Exam Help
<https://powcoder.com>
Add WeChat powcoder

Softmax regression

Classification by minimising cross-entropy loss

Srinandan Dasmahapatra

Classification: discrete output


Minimise deviation of prediction from annotation

- Given training set represented by points labelled **green** and **red**, variable Y ...

- ... where each point has two features

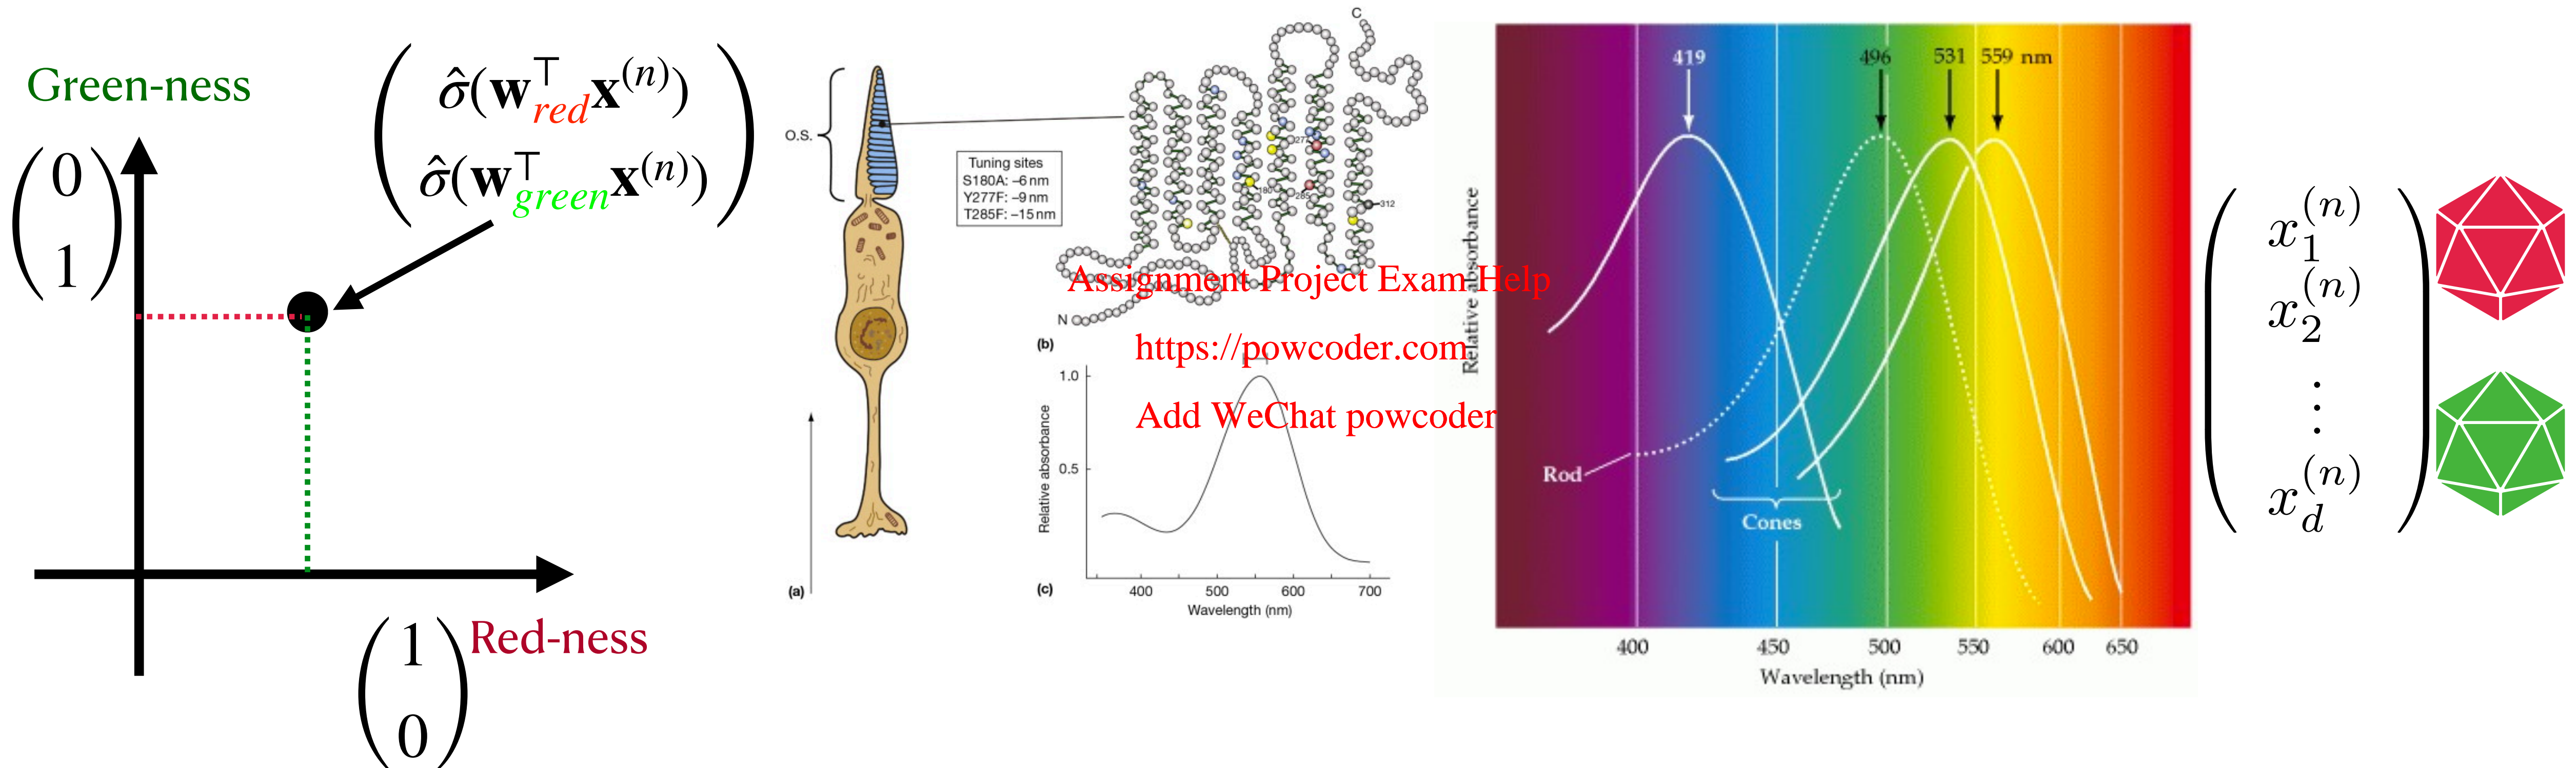
$$\mathcal{D} := \{((x_1^{(n)}, x_2^{(n)}), y^{(n)}) \mid n = 1, \dots, N\}$$

- Task: find function $f(x_1^{(n)}, x_2^{(n)}) = \hat{y}^{(n)}$ that reproduces given labels



Analogy with seeing in colour

Opsins (photopigments) in cones respond to colour preferentially



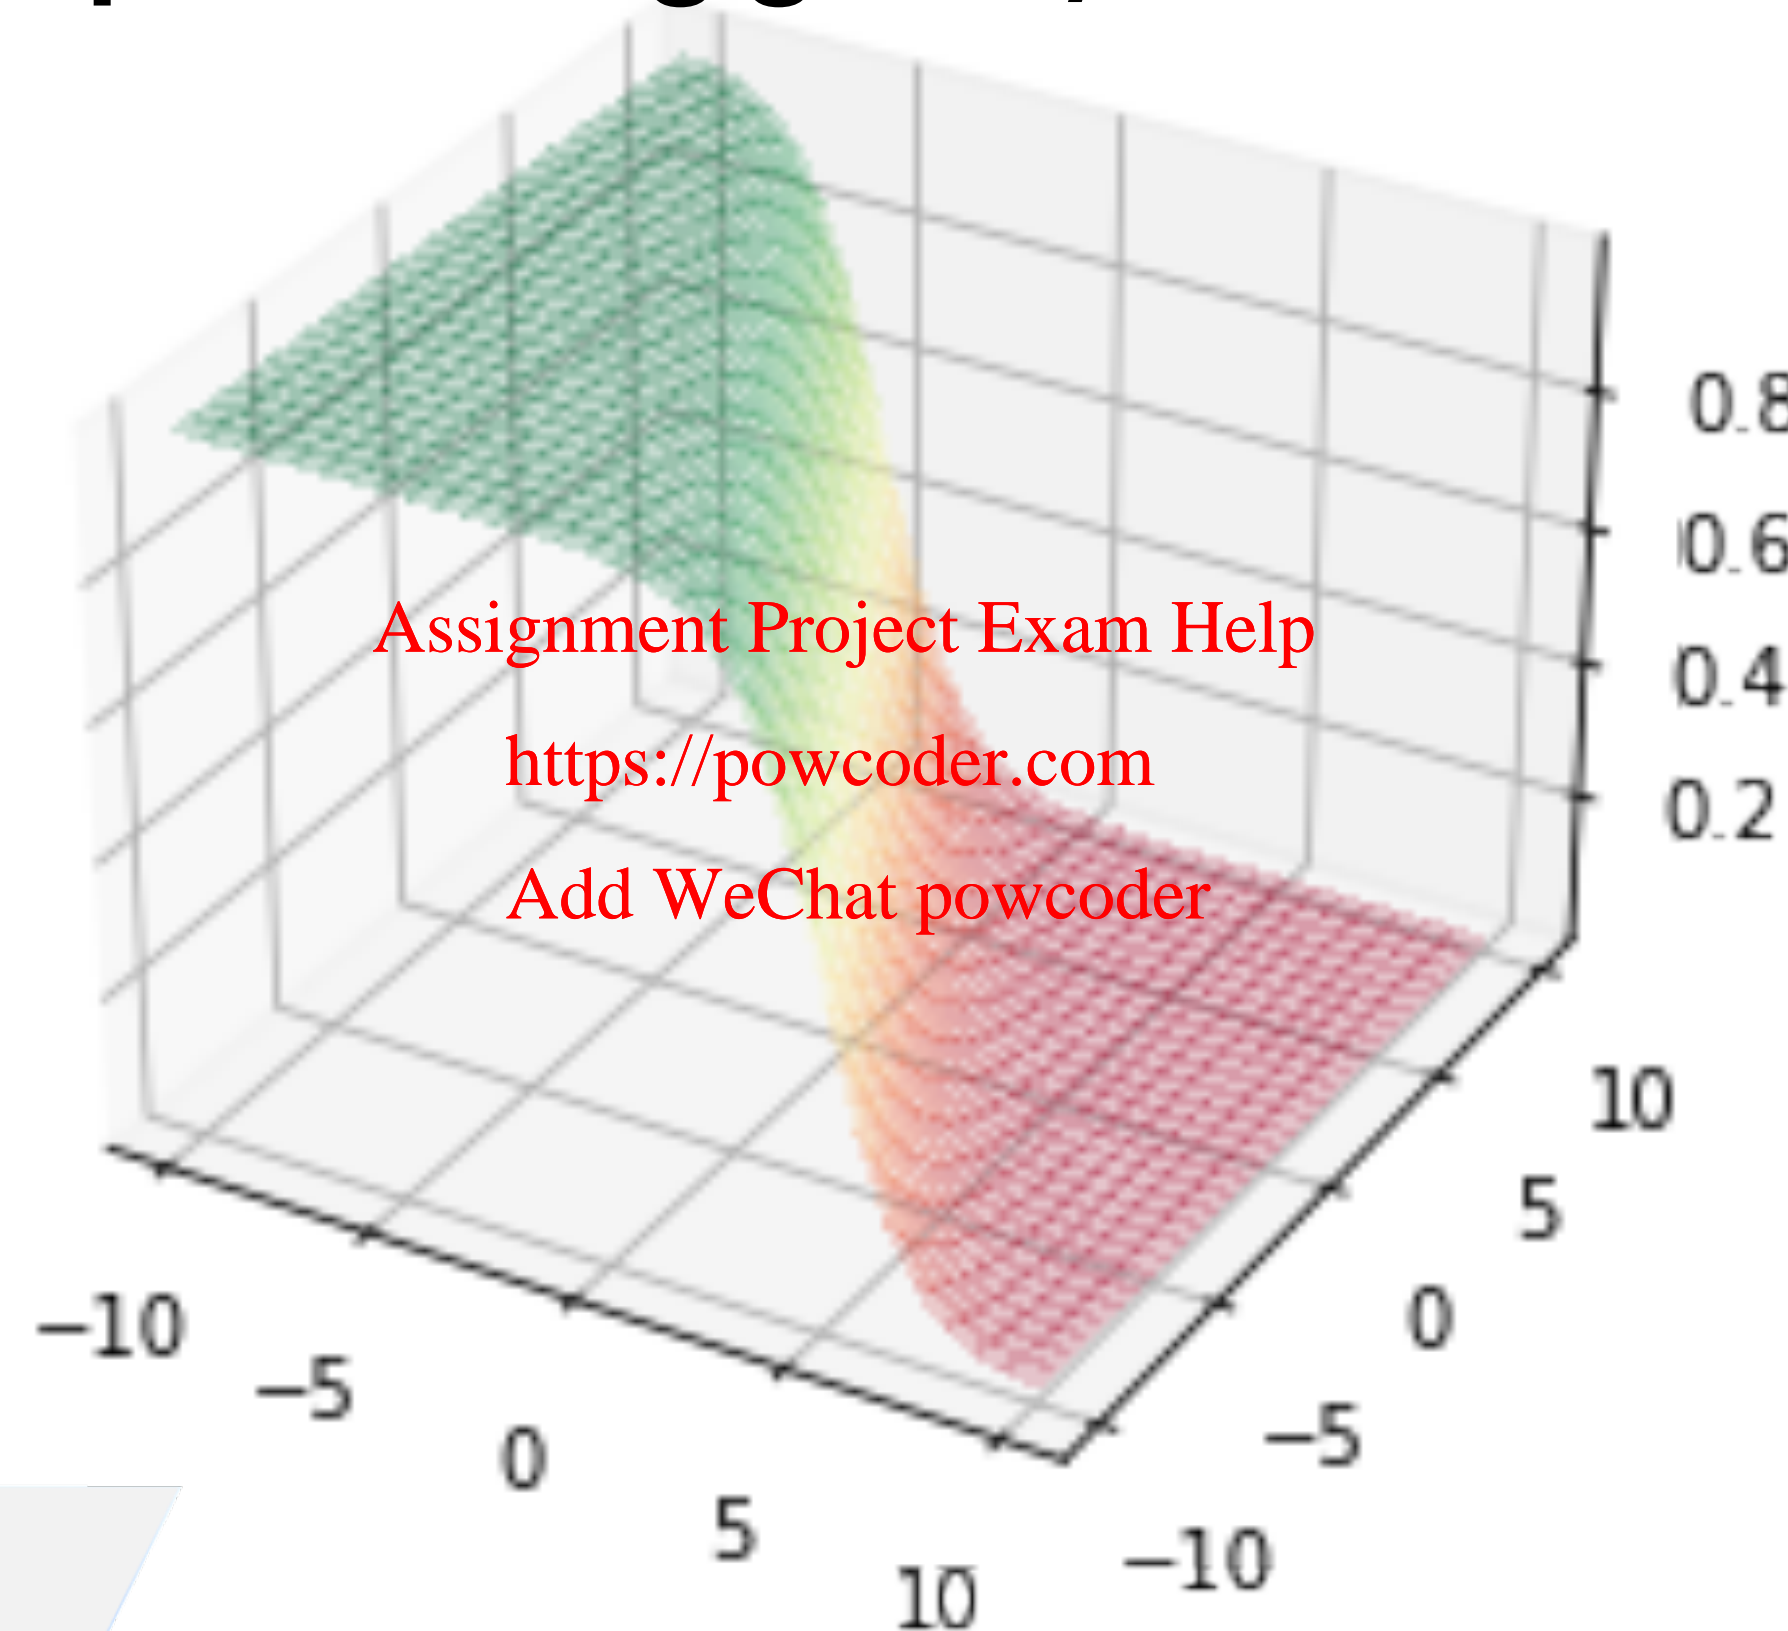

$$\mathbf{w}_{red} := (w_{red,1}, w_{red,2}, \dots, w_{red,d})$$

$$\mathbf{w}_{green} := (w_{green,1}, w_{green,2}, \dots, w_{green,d})$$

Find equation for decision boundary


Assign probability for each point being green/red

decision
boundary



$$f(x_1, x_2; \mathbf{w}) = w_0 + w_1x_1 + w_2x_2$$

$$\sigma(f) = \frac{1}{1 + \exp(-f)}$$




Learning = adjusting weights until agreement with data



Constructing scale for comparing predictions with training labels

- Output $f(\mathbf{x}^{(n)}; \mathbf{w}^i) \triangleq f_i^{(n)}, i = \text{red/green}, C = 2,$
- $0 \leq \sigma(f_i^{(n)}) \leq 1$ probability, with $\sum_{i=1}^C \sigma(f_i^{(n)}) = 1$
- Let $\hat{y}_i^{(n)} = \sigma(f_i^{(n)})$
- $y^{(n)} = \begin{pmatrix} 1 \\ 0 \end{pmatrix}$ or $\begin{pmatrix} 0 \\ 1 \end{pmatrix}$ for red/green
- Evaluation of classification: $\text{cost}(y^{(n)}, \hat{y}^{(n)})$
- Compare two different costs — quadratic and logarithmic
- Logarithm penalises mistakes more, also has a sharper drop (large gradient to guide weights to lower loss)



For a one component (scalar) output

Multiclass classification

Input: images 32 x 32 x 3 dimensions, Output: one-hot encodings: 10 dimensions

Here are the classes in the dataset, as well as 10 random images from each:

airplane



automobile




bird



cat




deer




dog




frog



horse



ship



truck



CIFAR-10: Example dataset for multi class classification

CAT $\mapsto \mathbf{e}_4 =$

$$\begin{pmatrix} 0 \\ 0 \\ 0 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{pmatrix}$$

SHIP $\mapsto \mathbf{e}_9 =$



$$\begin{pmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 1 \\ 0 \end{pmatrix}$$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

C-classes C different weight vectors \mathbf{w}^i



$$x \geq 0, y \geq 0, x + y = 1$$

- For each input vector (say a representation of an image or sound file), produce an output on the (C-1)-dimensional surface embedded in C-dimensional Euclidean space



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- Cost for input $\mathbf{x}^{(n)}$ = measure of mismatch between C-dimensional prediction $\hat{\sigma}(f(\mathbf{x}^{(n)}; \mathbf{w}^i))$ and true label \mathbf{e}_i

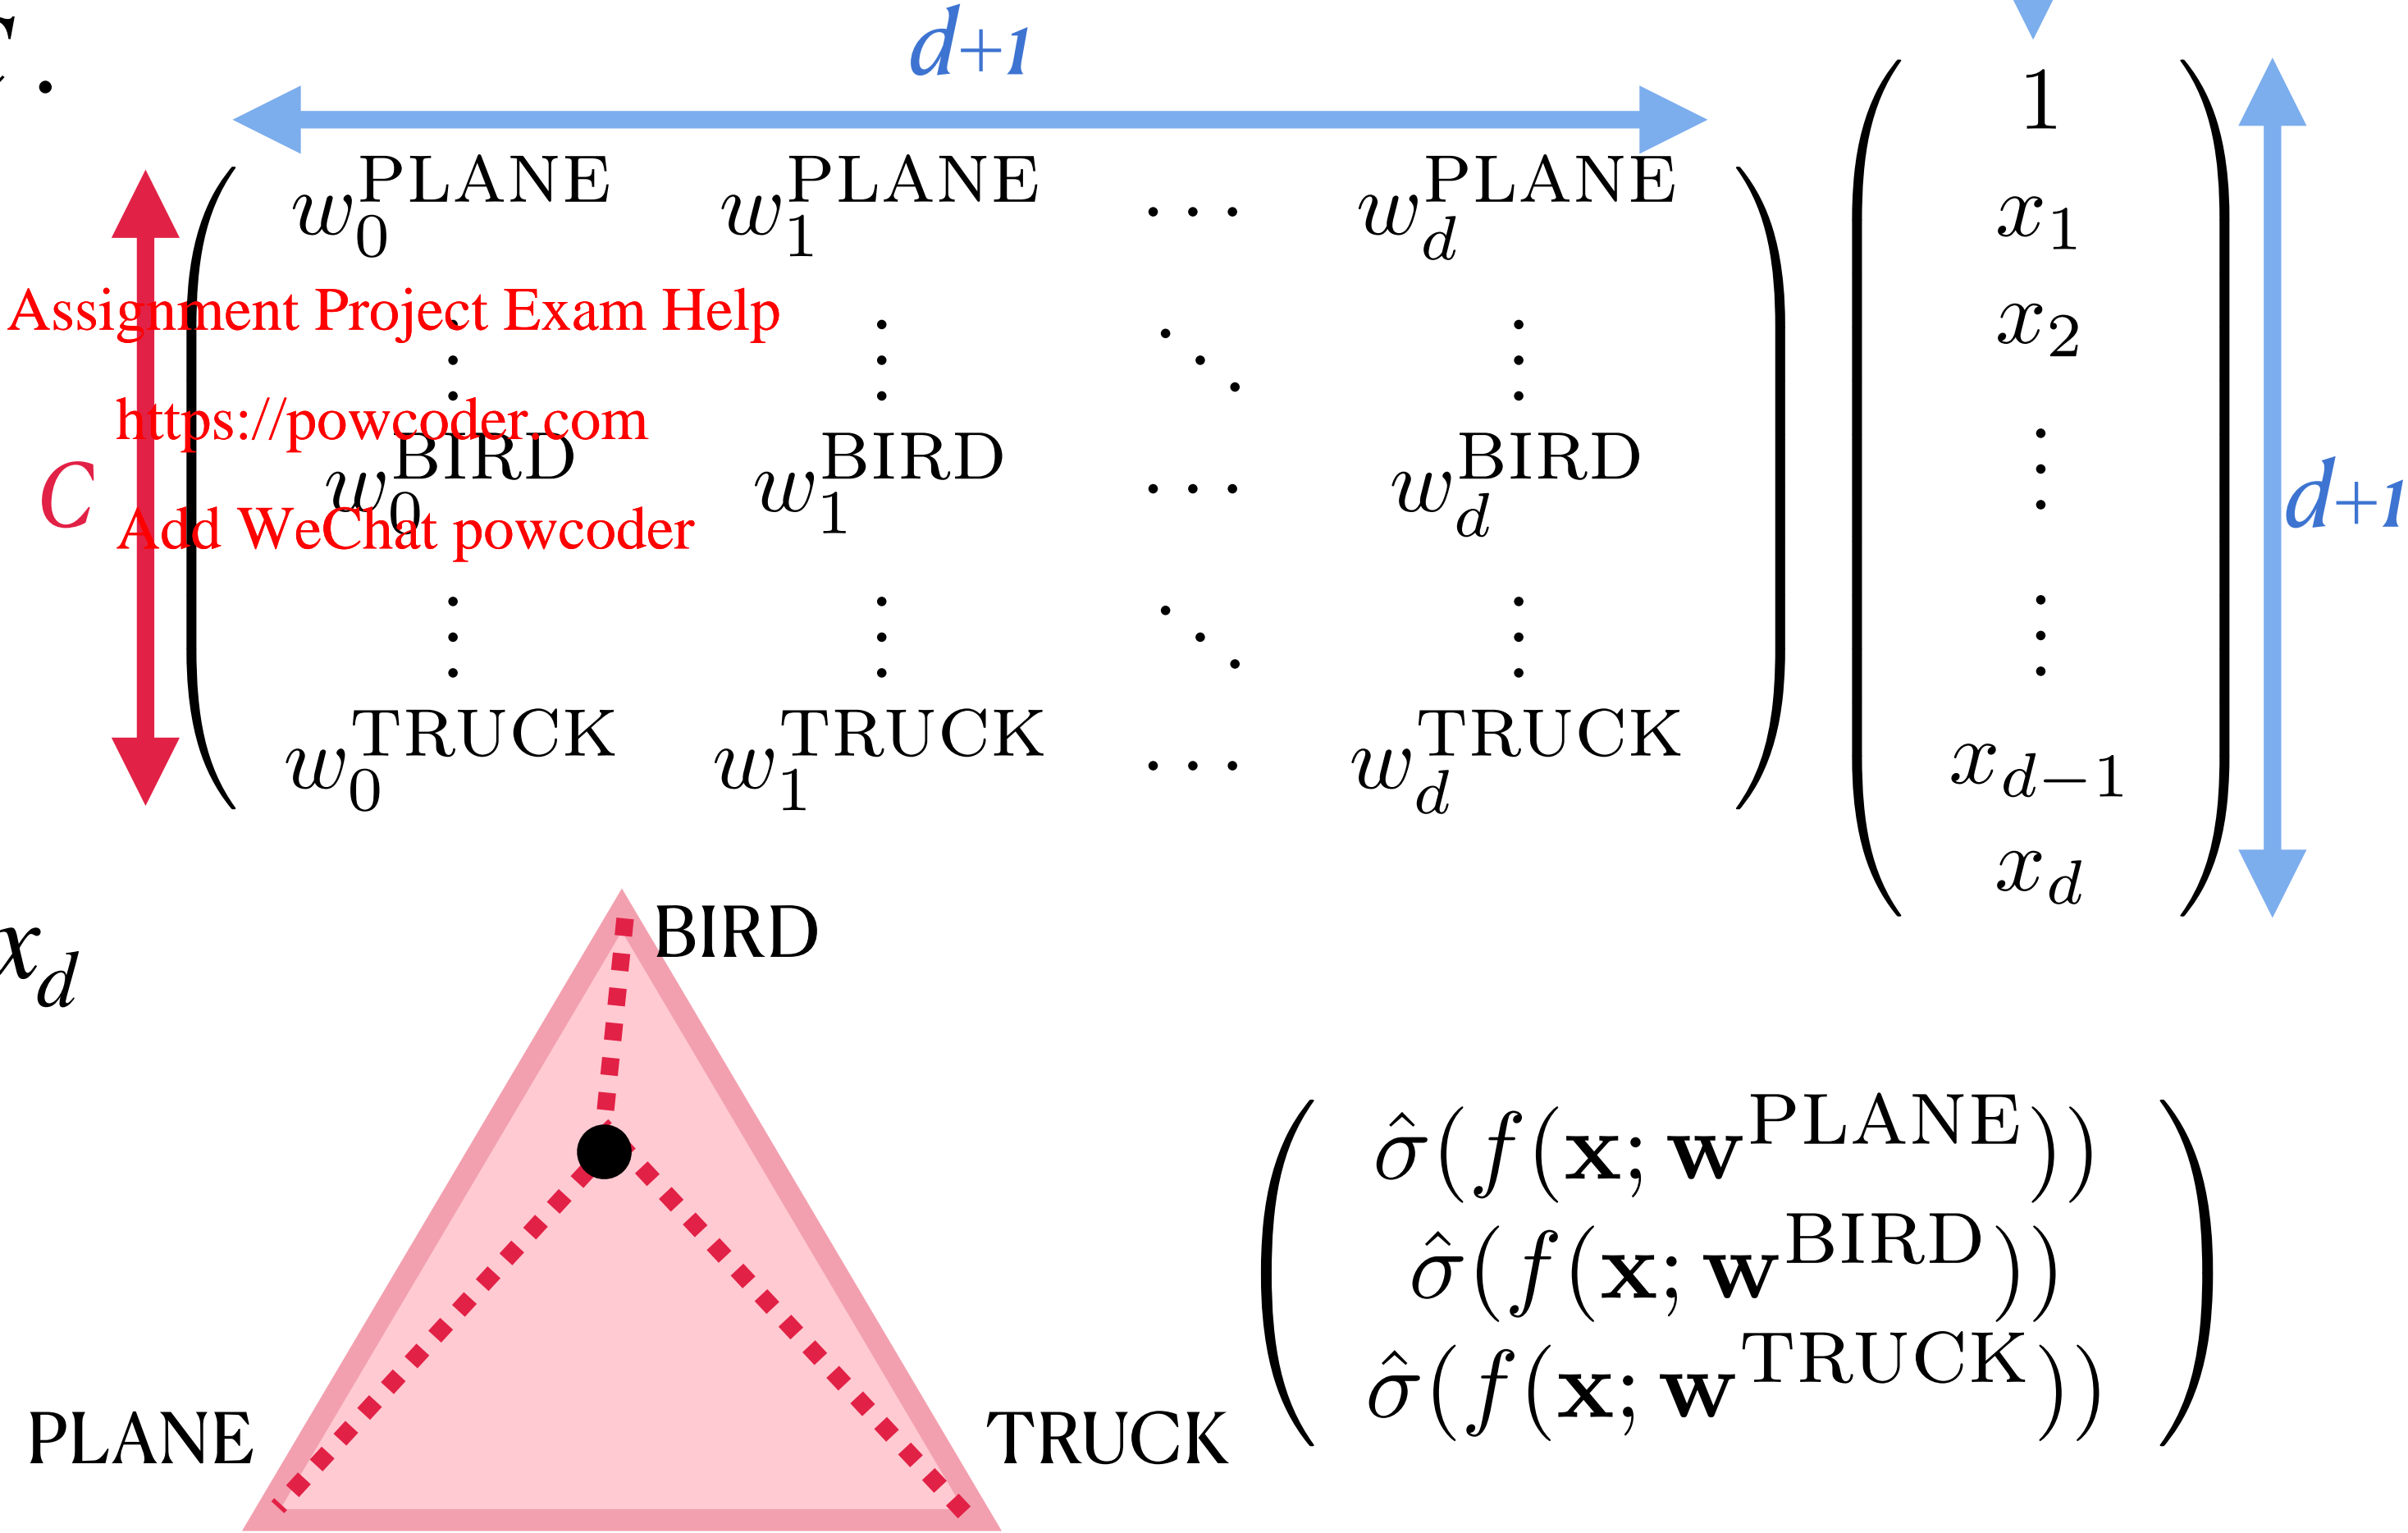
- Hat $\hat{\cdot}$ on $\hat{\sigma}$ indicates normalisation: entries add up to one: $\hat{\sigma}(f(\mathbf{x}; \mathbf{w}^1)) + \hat{\sigma}(f(\mathbf{x}; \mathbf{w}^2)) + \hat{\sigma}(f(\mathbf{x}; \mathbf{w}^3)) = 1$



Multiclass classification

Weight vectors for each class

- $\mathbf{w}^c = (w_0^c, w_1^c, \dots, w_d^c), c = 1, \dots, C.$
- C - number of classes, 10 for CIFAR-10
- d - dimensionality of data,
 $\mathbf{x} = (x_1, x_2, \dots, x_d)$
- $f(\mathbf{x}; \mathbf{w}^c) = w_0^c \cdot 1 + w_1^c x_1 + \dots + w_d^c x_d$
 : for each input data point,
 compute output for all classes



Set up gradient descent of loss for classification

Re-phrasing what has been done

- For each class each data point $\mathbf{x}^{(n)}$ is assigned a score $s_c^{(n)} = f(\mathbf{x}^{(n)}; \mathbf{w}^c)$, $c = 1, \dots, C$
- Choose the largest of the C scores as the predicted class for $\mathbf{x}^{(n)}$
 - $c^* = \arg \max_{c \in \{1, \dots, C\}} s_c^{(n)}$
- Replace max by softmax: $\max(s_1, s_2, s_3) \longrightarrow \text{softmax}(s_1, s_2, s_3) = \ln(e^{s_1} + e^{s_2} + e^{s_3})$
- Exponential function: monotonic in argument ($x \nearrow \implies e^x \nearrow$)
- Normalise exponential scores: $s_c^{(n)} \mapsto \frac{e^{s_c^{(n)}}}{e^{s_1^{(n)}} + e^{s_2^{(n)}} + \dots + e^{s_C^{(n)}}} =: \hat{\sigma}(s_c^{(n)}) = [\hat{y}^{(n)}]_c$
- Treat component c of $[\hat{y}^{(n)}]_c = \hat{\sigma}(s_c^{(n)})$ as probability that $\mathbf{x}^{(n)}$ belongs to class c : $P(c | \mathbf{x}^{(n)})$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Multi-class loss function: cross entropy

Measures information about label distribution from input data and choice of weights

- For each data point $\mathbf{x}^{(n)}$ sum over costs $-\sum_{c=1}^C y_c^{(n)} \ln \hat{y}_c^{(n)}$ for all classes
- Sum costs over all data points $L(\mathbf{W}) := L(\{\mathbf{w}^1, \dots, \mathbf{w}^C\}) = -\sum_{n=1}^N \sum_{c=1}^C y_c^{(n)} \ln \hat{y}_c^{(n)}$, called cross-entropy.
Assignment Project Exam Help
<https://powcoder.com>
Add WeChat powcoder
- **Eg:** target $y^{(n)} = \begin{pmatrix} 0 \\ 1 \\ 0 \\ 0 \end{pmatrix}$ prediction $\hat{y}^{(n)} = \begin{pmatrix} \hat{y}_1^{(n)} \\ \hat{y}_2^{(n)} \\ \hat{y}_3^{(n)} \\ \hat{y}_4^{(n)} \end{pmatrix} : -(0 \cdot \ln \hat{y}_1^{(n)} + 1 \cdot \ln \hat{y}_2^{(n)} + 0 \cdot \ln \hat{y}_3^{(n)} + 0 \cdot \ln \hat{y}_4^{(n)}) = -\ln \hat{y}_2^{(n)}$
- $L(\mathbf{W}) = -\ln \left(\hat{y}_{c_1}^{(1)} \cdot \hat{y}_{c_2}^{(2)} \cdots \hat{y}_{c_N}^{(N)} \right) = -\sum_{n=1}^N \ln \hat{y}_{c_n}^{(n)}$: reduce negative of log(predicted probabilities)

Gradient descent on cross-entropy finds optimal weights


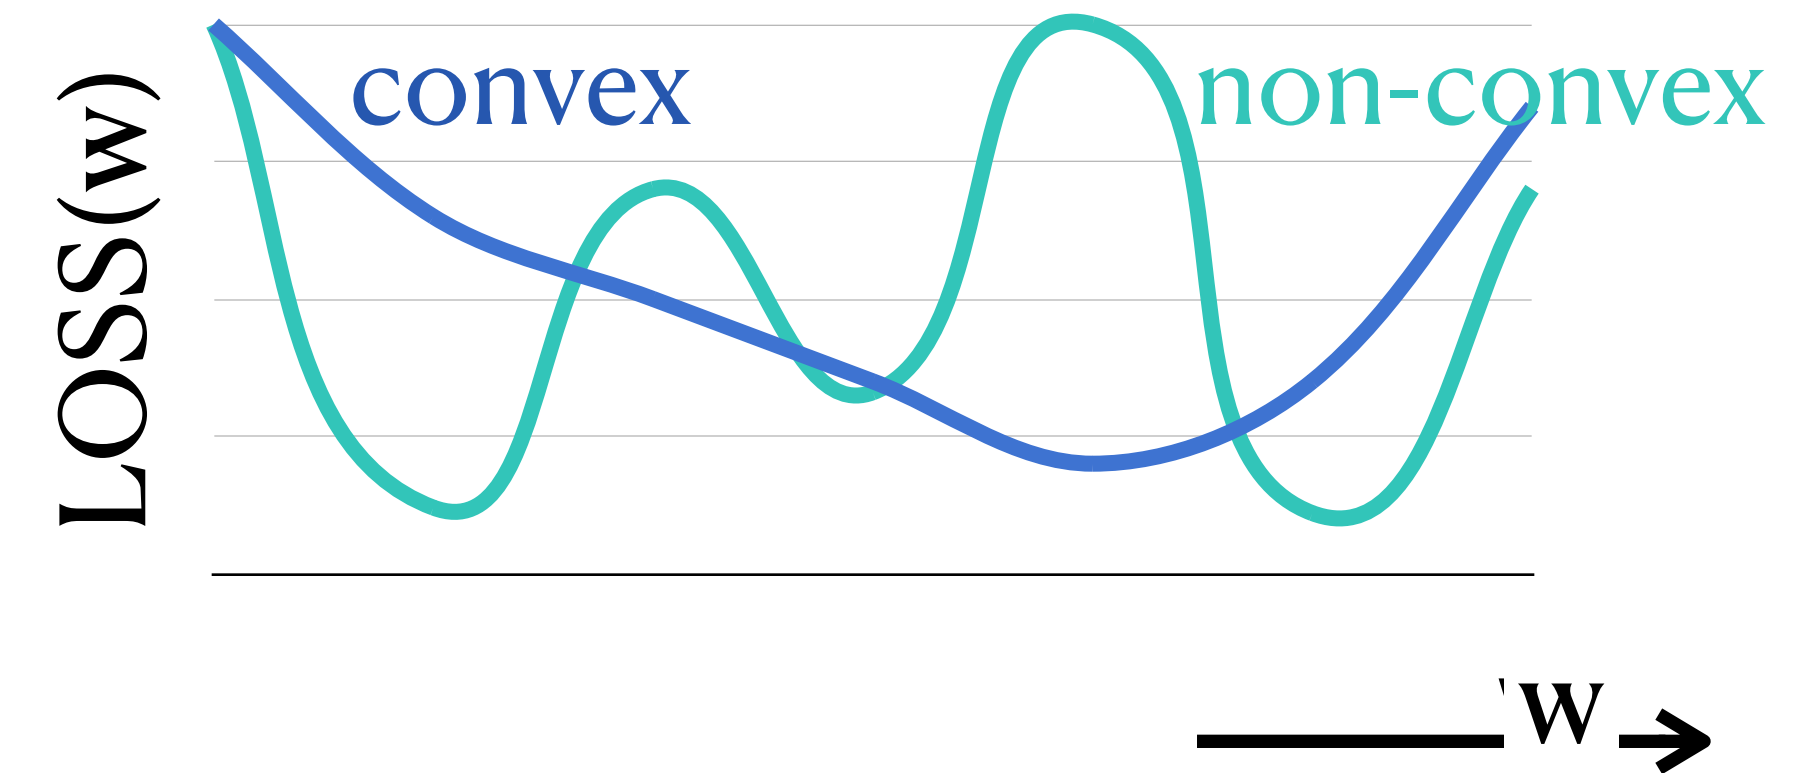
For linear maps f , cross-entropy is convex

- Learning: Reduce $L(\mathbf{W})$ by changing weights
- $\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \eta \nabla_{\mathbf{w}} L(\mathbf{W})$
- **All** weights are contained in \mathbf{w}
- Jupyter notebook

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder




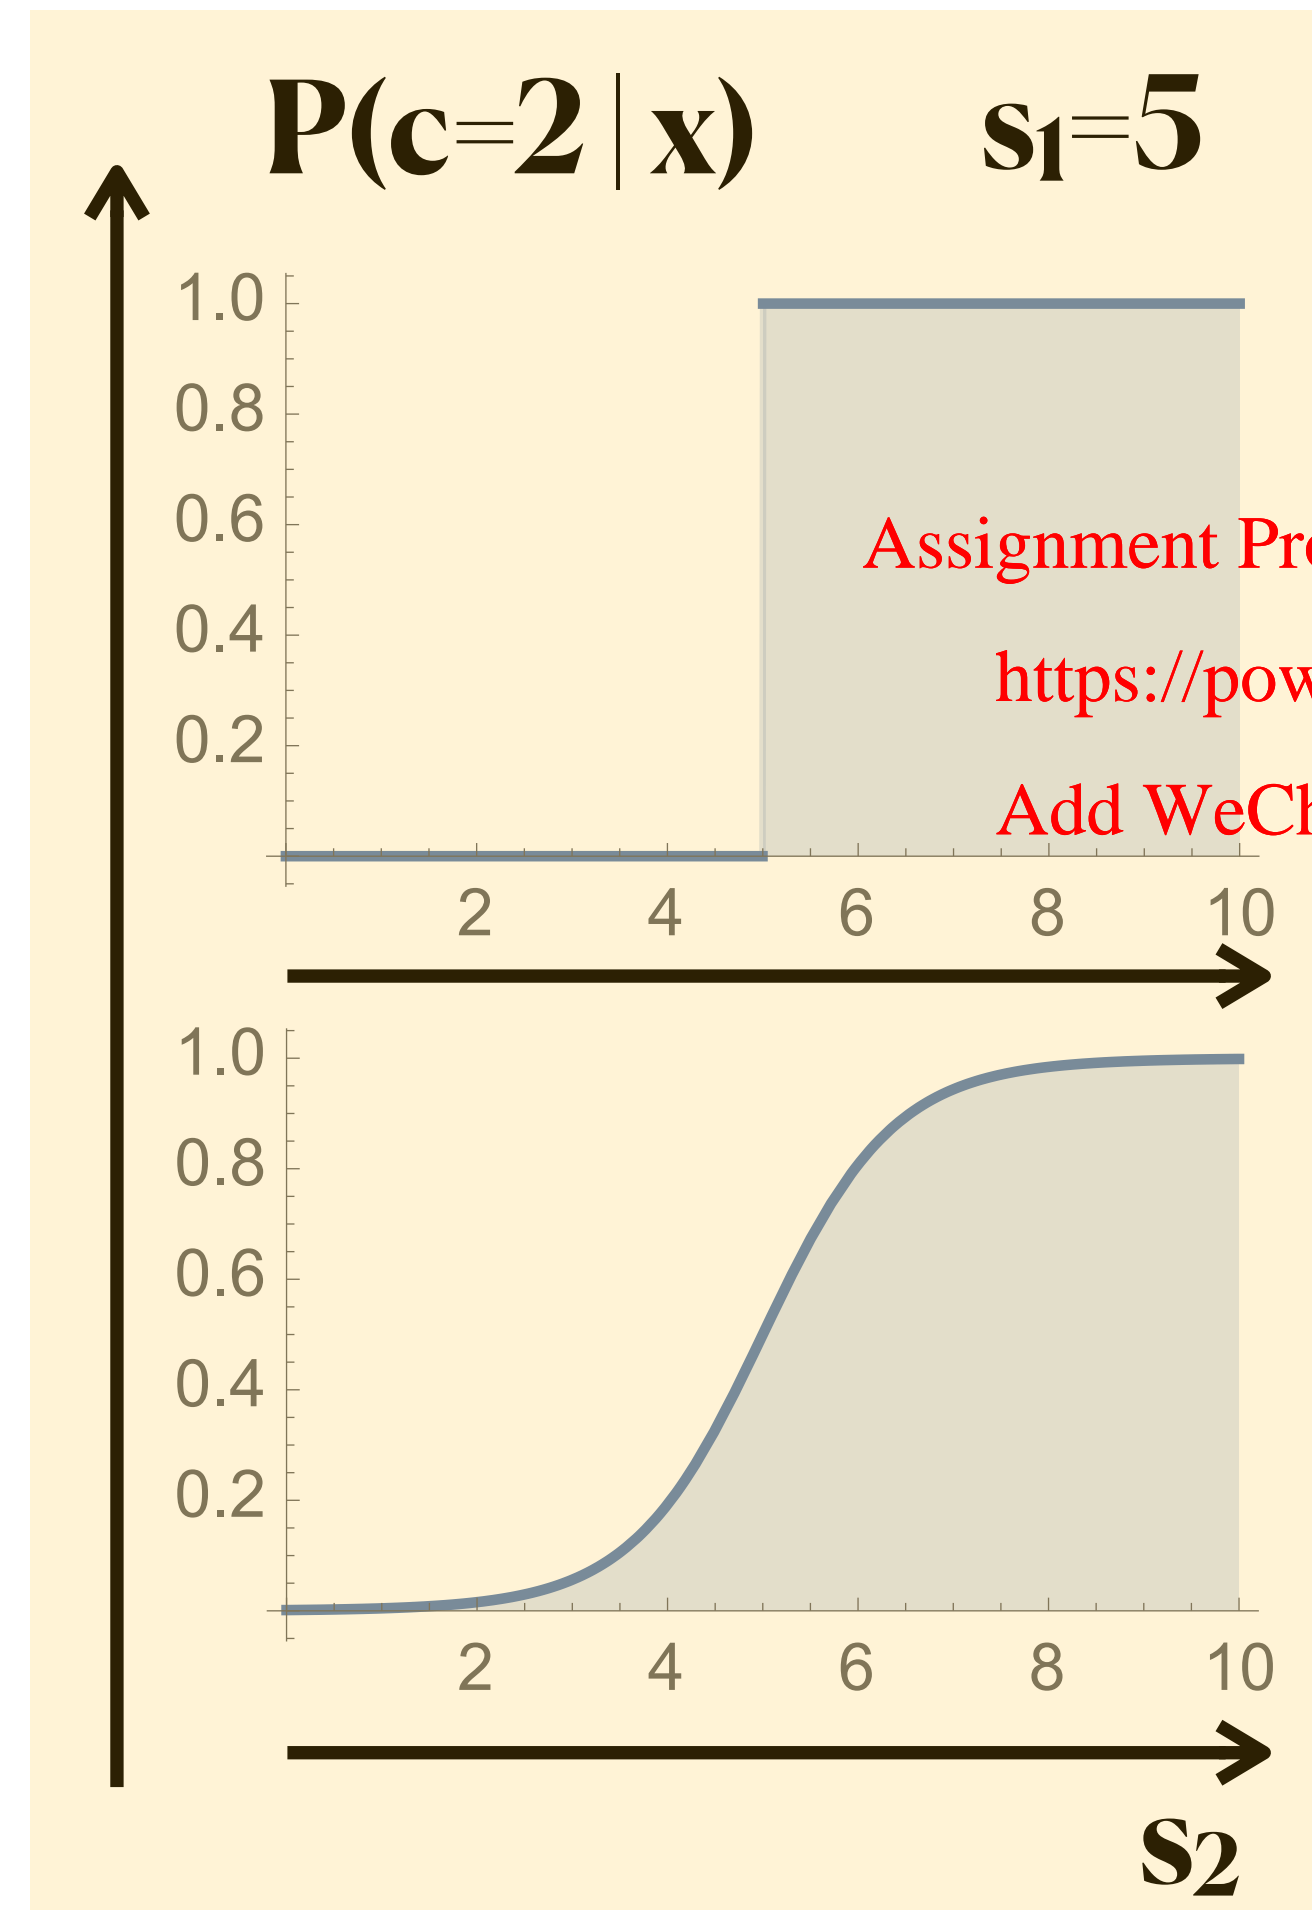
Example: data x ; 2-class problem

Compare probability assignment for arg max with arg softmax

$$\frac{f[0, (s-5)]}{f[0, (s-5)] + f[0, -(s-5)]}$$

- $f=\text{Max}$

- $f=\text{Softmax}$



Assignment Project Exam Help
Lab 2
<https://powcoder.com>
Add WeChat powcoder